

LSF-CE/SLURM-CodeEngine-Gateway

Christof Westhues

IBM

cwesthues@de.ibm.com

2024/10/24



IBM



What is it and why?

A gateway to run LSF-CE or SLURM workload (,jobs‘) on (dynamic) IBM CodeEngine instances

Run workload on ,**serverless computenodes**‘

No complex handling of (VPC) compute nodes, network, storage

For simple, basic workload

Value proposition

- ✓ **dynamically scale compute resources** on-demand **for peak workloads**
- ✓ improve **cost efficiency** through Pay-Per-Use Model
- ✓ **reduce operational overhead** ("No Infrastructure Management")
- ✓ **secure** containerized job environments
- ✓ **easy Job Submission** and Management
- ✓ **enhanced Job Distribution** and **Fault Tolerance**
- ✓ stateful **handling of data** through COS mounts to container
- ✓ **easy to implement** (and demo)

Potential applications

EDA:

Ansys HFSS

Cadence Xcelium

Keysight ADS

Mentor Graphics Calibre

Silvaco TCAD

Synopsys Sentaurus

Synopsys VCS

Manufacturing:

ABAQUS *

Altair HyperWorks *

ANSYS Additive

ANSYS Fluent *

ANSYS Mechanical

ANSYS Multiphysics

AutoForm

COMSOL Multiphysics *

OpenFOAM *

Star-CCM+ *

Oil & Gas:

ECLIPSE (by
Schlumberger) *

Schlumberger Petrel

Schlumberger VISAGE *

Lifesciences:

AlphaFold *

AMBER *

AutoDock

BLAST

DOCK

GATK

GROMACS *

NAMD *

ROSETTA

Others:

Matlab

R

Tensorflow

Octave

PyTorch

* typically uses MPI/Infiniband

Install your SW/apps (example) in docker container:

```
cat > Dockerfile <<EOF
FROM docker.io/library/ubuntu:latest
RUN apt-get -yq update ; \
    apt -y install curl git python3-pip unzip wget
COPY my_app /usr/bin/my_app
EOF

docker build -t docker.io/${DOCKERUSER}/ubuntu .

docker push docker.io/${DOCKERUSER}/ubuntu
```

Access to data through COS

What

Where

```
s3fs ${COSBUCKET} ${MOUNTPOINT} \  
-o url=https://s3.eu-de.cloud-object-storage.appdomain.cloud  
mkdir -p ${MOUNTPOINT}/${JOBID}  
cp INPUT.txt ${MOUNTPOINT}/${JOBID}  
umount ${MOUNTPOINT}
```

1.

Mount a COS bucket
and copy input data

submitting/
master
node

```
cat /mnt/${JOBID}/INPUT.txt | ... > /mnt/${JOBID}/OUTPUT.txt
```

2.

Inside the container,
„the job“

inside container

```
s3fs ${COSBUCKET} ${MOUNTPOINT} \  
-o url=https://s3.eu-de.cloud-object-storage.appdomain.cloud  
cp ${MOUNTPOINT}/${JOBID}/OUTPUT.txt ...  
umount ${MOUNTPOINT}
```

3.

Mount a COS bucket
and copy output data

submitting/
master
node

Jumping straight into tech stuff...

Begin Application

```
NAME                = codeengine

RES_REQ             = select[codeengine

CONTAINER           = docker[image(${LSB_CODEENGINE_IMAGE}) \
                        options(${LSB_CODEENGINE_OPTIONS})]

DESCRIPTION         = Codeengine User Service

EXEC_DRIVER         = context[user(default)] \
                    starter[codeengine-starter.sh] \
                    controller[codeengine-control.sh]

PRE_EXEC            = codeengine-preexec.sh

POST_EXEC           = codeengine-postexec.sh

JOB__INCLUDE__POSTPROC = Y
```

End Application

Some more details on codeengine-starter.sh (simplified)

```
#!/bin/sh

ibmcloud login ...

ibmcloud ce project select -n ...

ibmcloud ce job create --name ... --image ... --argument ${CMD}

add-volume-mount-to-job.sh ${LSB_JOBID} s3fstest2 /mnt

ibmcloud ce jobrun submit --job ${LSB_JOBID}

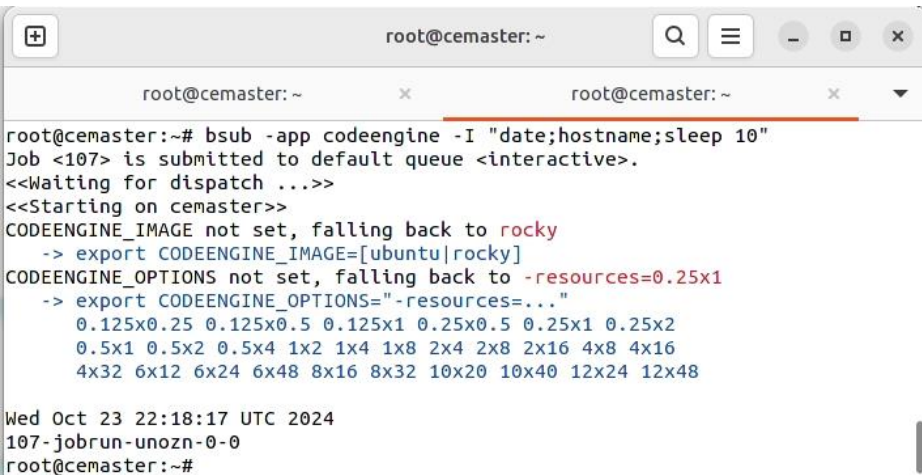
ibmcloud ce jobrun logs -f -r --jobrun ${JOBRUN_NAME}

ibmcloud ce jobrun delete --name ${JOBRUN_NAME} -f

ibmcloud ce job delete --name ${LSB_JOBID} -f
```

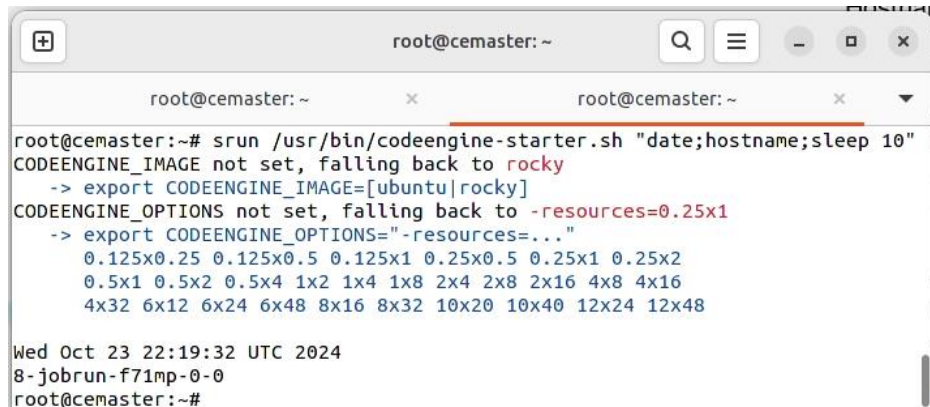

Running LSF-CE and SLURM interactive jobs on CodeEngine

LSF-CE



```
root@cemaster: ~  
root@cemaster: ~  
root@cemaster:~# bsub -app codeengine -I "date;hostname;sleep 10"  
Job <107> is submitted to default queue <interactive>.  
<<Waiting for dispatch ...>>  
<<Starting on cemaster>>  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
  
Wed Oct 23 22:18:17 UTC 2024  
107-jobrun-unozn-0-0  
root@cemaster:~#
```

SLURM



```
root@cemaster: ~  
root@cemaster: ~  
root@cemaster:~# srun /usr/bin/codeengine-starter.sh "date;hostname;sleep 10"  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
  
Wed Oct 23 22:19:32 UTC 2024  
8-jobrun-f71mp-0-0  
root@cemaster:~#
```

Running LSF-CE and SLURM interactive jobs on CodeEngine

LSF-CE

```
root@cemaster: ~  
root@cemaster: ~  
root@cemaster:~# bsub -app codeengine -I "cat /etc/os-release"  
Job <108> is submitted to default queue <interactive>.  
<<Waiting for dispatch ...>>  
<<Starting on cemaster>>  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
  
NAME="Rocky Linux"  
VERSION="8.6 (Green Obsidian)"  
ID="rocky"  
ID_LIKE="rhel centos fedora"  
VERSION_ID="8.6"  
PLATFORM_ID="platform:el8"  
PRETTY_NAME="Rocky Linux 8.6 (Green Obsidian)"  
ANSI_COLOR="0;32"  
CPE_NAME="cpe:/o:rocky:rocky:8:GA"  
HOME_URL="https://rockylinux.org/"  
BUG_REPORT_URL="https://bugs.rockylinux.org/"  
ROCKY_SUPPORT_PRODUCT="Rocky Linux"  
ROCKY_SUPPORT_PRODUCT_VERSION="8"  
REDHAT_SUPPORT_PRODUCT="Rocky Linux"  
REDHAT_SUPPORT_PRODUCT_VERSION="8"  
root@cemaster:~#
```

SLURM

```
root@cemaster: ~  
root@cemaster: ~  
root@cemaster:~# srun /usr/bin/codeengine-starter.sh "cat /etc/os-release"  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
  
NAME="Rocky Linux"  
VERSION="8.6 (Green Obsidian)"  
ID="rocky"  
ID_LIKE="rhel centos fedora"  
VERSION_ID="8.6"  
PLATFORM_ID="platform:el8"  
PRETTY_NAME="Rocky Linux 8.6 (Green Obsidian)"  
ANSI_COLOR="0;32"  
CPE_NAME="cpe:/o:rocky:rocky:8:GA"  
HOME_URL="https://rockylinux.org/"  
BUG_REPORT_URL="https://bugs.rockylinux.org/"  
ROCKY_SUPPORT_PRODUCT="Rocky Linux"  
ROCKY_SUPPORT_PRODUCT_VERSION="8"  
REDHAT_SUPPORT_PRODUCT="Rocky Linux"  
REDHAT_SUPPORT_PRODUCT_VERSION="8"  
root@cemaster:~#
```

Running LSF-CE and SLURM interactive jobs on CodeEngine

LSF-CE

```
root@cemaster: ~  
root@cemaster: ~# bsub < /tmp/lsf_teststdout.sh  
Job <104> is submitted to default queue <interactive>.  
<<Waiting for dispatch ...>>  
<<Starting on cemaster>>  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
  
Run 5 iterations  
Run 1  
Thu Oct 24 08:03:10 UTC 2024  
Run 2  
Thu Oct 24 08:03:15 UTC 2024  
Run 3  
Thu Oct 24 08:03:20 UTC 2024  
Run 4  
Thu Oct 24 08:03:25 UTC 2024  
Run 5  
Thu Oct 24 08:03:30 UTC 2024  
root@cemaster:~#
```

SLURM

```
root@cemaster: ~  
root@cemaster:~# srun /usr/bin/codeengine-starter.sh /tmp/slurm_teststdout.sh  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
  
Run 5 iterations  
Run 1  
Thu Oct 24 08:04:50 UTC 2024  
Run 2  
Thu Oct 24 08:04:55 UTC 2024  
Run 3  
Thu Oct 24 08:05:00 UTC 2024  
Run 4  
Thu Oct 24 08:05:05 UTC 2024  
Run 5  
Thu Oct 24 08:05:10 UTC 2024  
root@cemaster:~#
```

Running LSF-CE and SLURM batch jobs on CodeEngine

LSF-CE

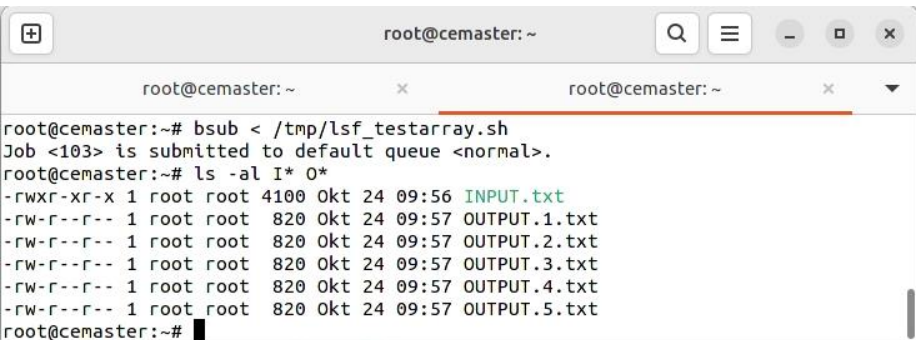
```
root@cemaster: ~  
root@cemaster:~# bsub < /tmp/lsf_testjob.sh  
Job <102> is submitted to default queue <interactive>.  
<<Waiting for dispatch ...>>  
<<Starting on cemaster>>  
CODEENGINE_IMAGE not set, falling back to rocky  
-> export CODEENGINE_IMAGE=[ubuntu|rocky]  
CODEENGINE_OPTIONS not set, falling back to -resources=0.25x1  
-> export CODEENGINE_OPTIONS="-resources=..."  
0.125x0.25 0.125x0.5 0.125x1 0.25x0.5 0.25x1 0.25x2  
0.5x1 0.5x2 0.5x4 1x2 1x4 1x8 2x4 2x8 2x16 4x8 4x16  
4x32 6x12 6x24 6x48 8x16 8x32 10x20 10x40 12x24 12x48  
#####  
cat /mnt/lsfce-102/INPUT.txt | sort --ignore-case > /mnt/lsfce-102/OUTPUT.txt:  
#####  
root@cemaster:~# ls -al I* 0*  
-rwxr-xr-x 1 root root 4100 0kt 24 09:49 INPUT.txt  
-rw-r--r-- 1 root root 4100 0kt 24 09:50 OUTPUT.txt  
root@cemaster:~#
```

SLURM

```
root@cemaster: ~  
root@cemaster:~# sbatch /tmp/slurm_testjob.sh  
Submitted batch job 2  
root@cemaster:~# ls -al I* 0*  
-rwxr-xr-x 1 root root 4100 0kt 24 09:51 INPUT.txt  
-rw-r--r-- 1 root root 4100 0kt 24 09:51 OUTPUT.txt  
root@cemaster:~#
```

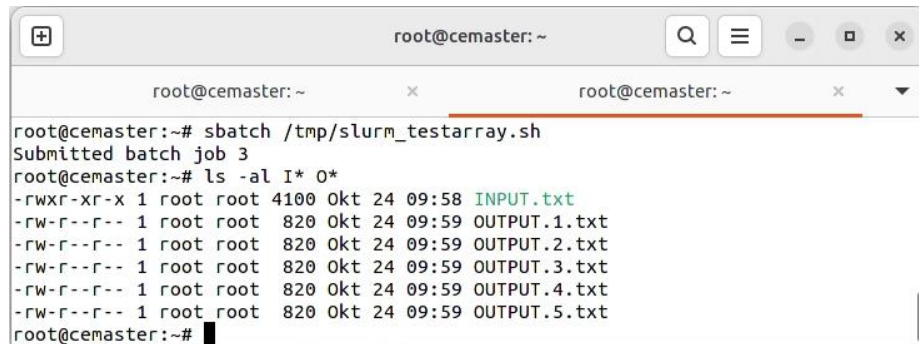
Running LSF-CE and SLURM array jobs on CodeEngine

LSF-CE



```
root@cemaster: ~  
root@cemaster:~# bsub < /tmp/lsf_testarray.sh  
Job <103> is submitted to default queue <normal>.  
root@cemaster:~# ls -al I* 0*  
-rwxr-xr-x 1 root root 4100 Okt 24 09:56 INPUT.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:57 OUTPUT.1.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:57 OUTPUT.2.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:57 OUTPUT.3.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:57 OUTPUT.4.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:57 OUTPUT.5.txt  
root@cemaster:~#
```

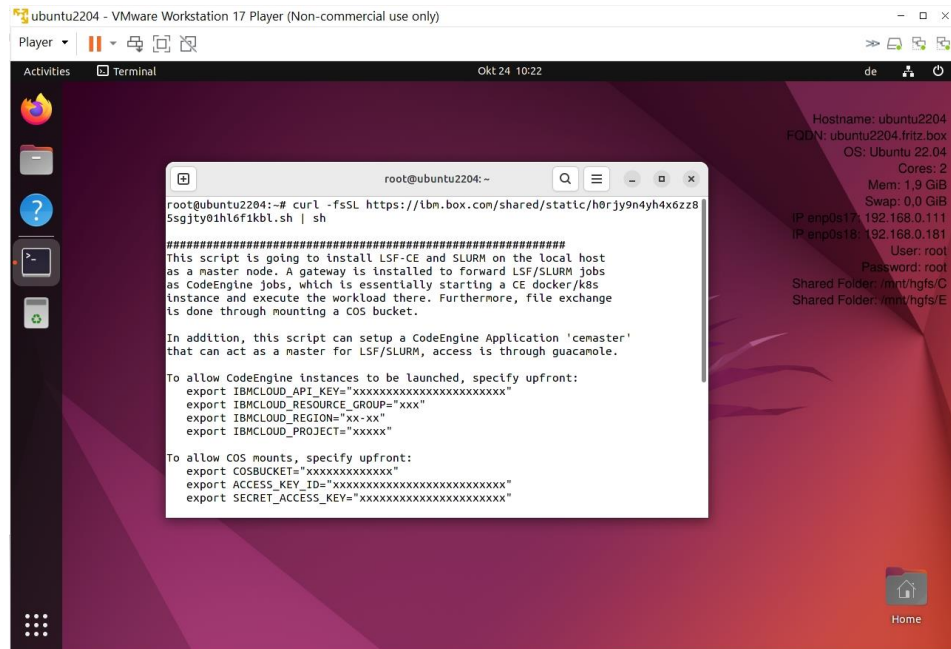
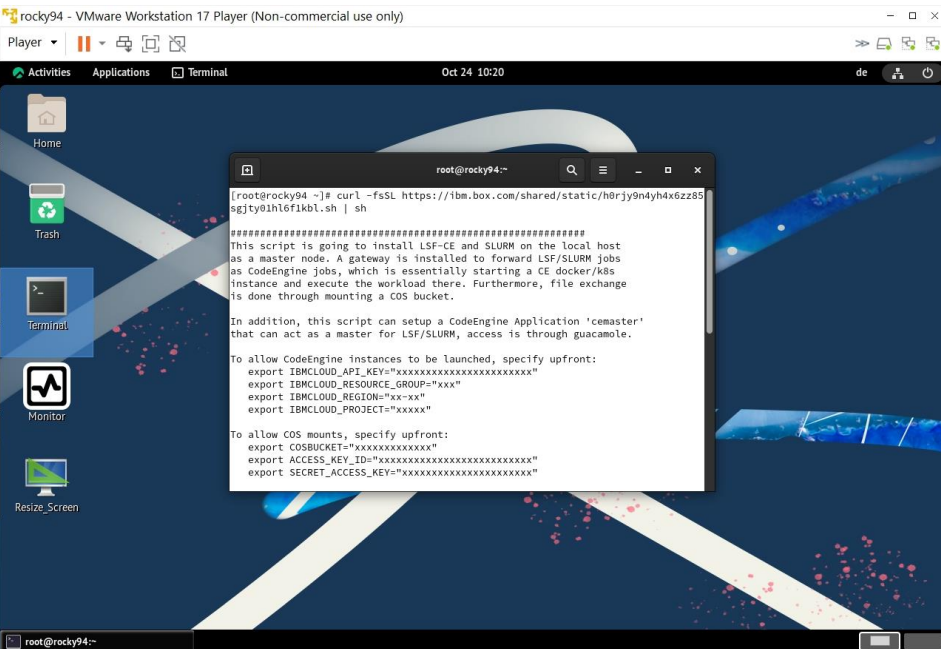
SLURM



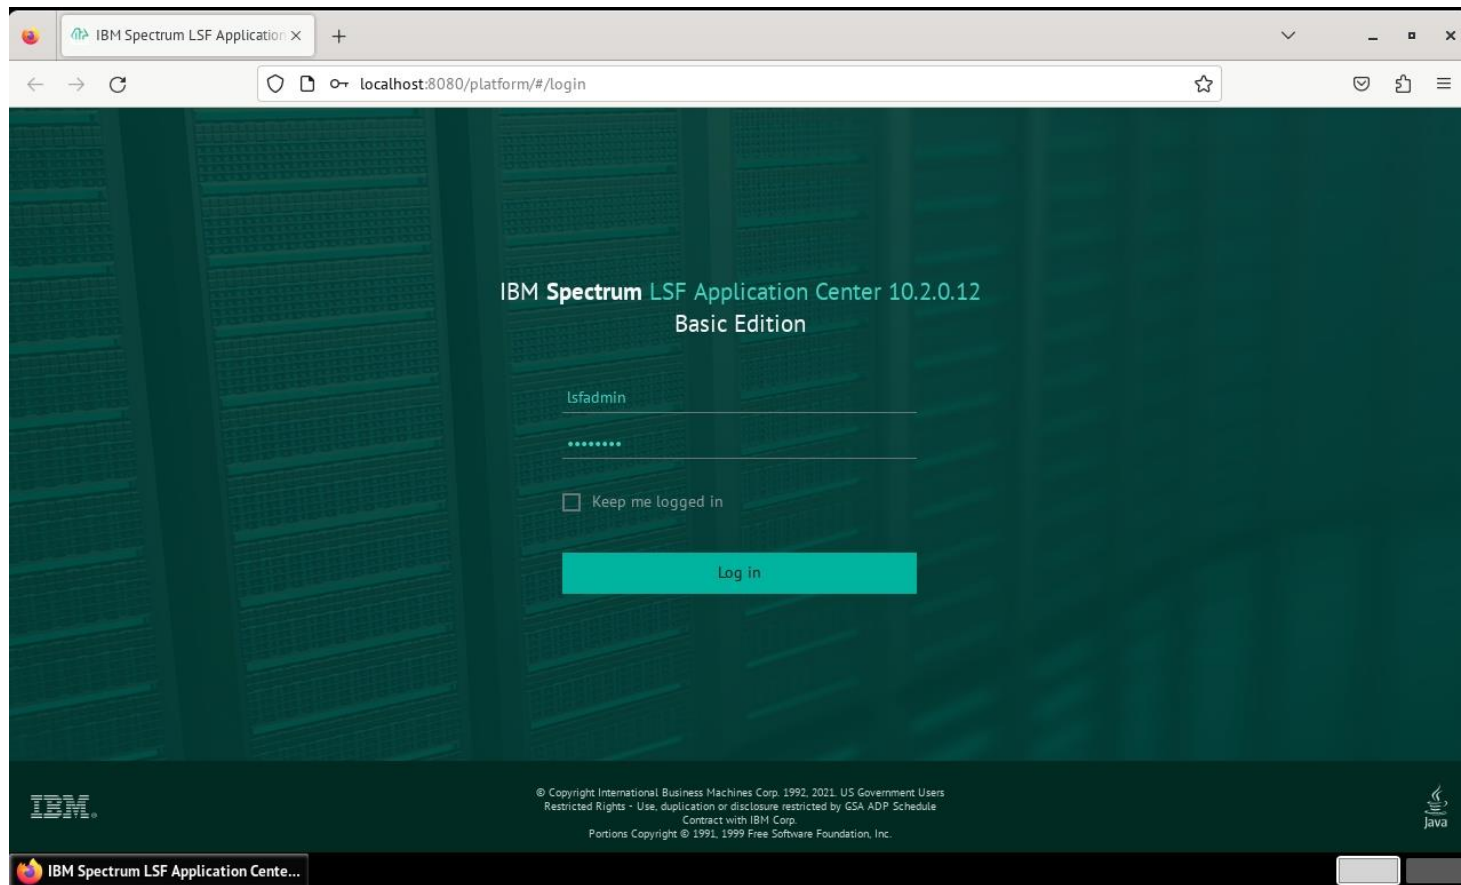
```
root@cemaster: ~  
root@cemaster:~# sbatch /tmp/slurm_testarray.sh  
Submitted batch job 3  
root@cemaster:~# ls -al I* 0*  
-rwxr-xr-x 1 root root 4100 Okt 24 09:58 INPUT.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:59 OUTPUT.1.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:59 OUTPUT.2.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:59 OUTPUT.3.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:59 OUTPUT.4.txt  
-rw-r--r-- 1 root root 820 Okt 24 09:59 OUTPUT.5.txt  
root@cemaster:~#
```


Install on Rocky and Ubuntu

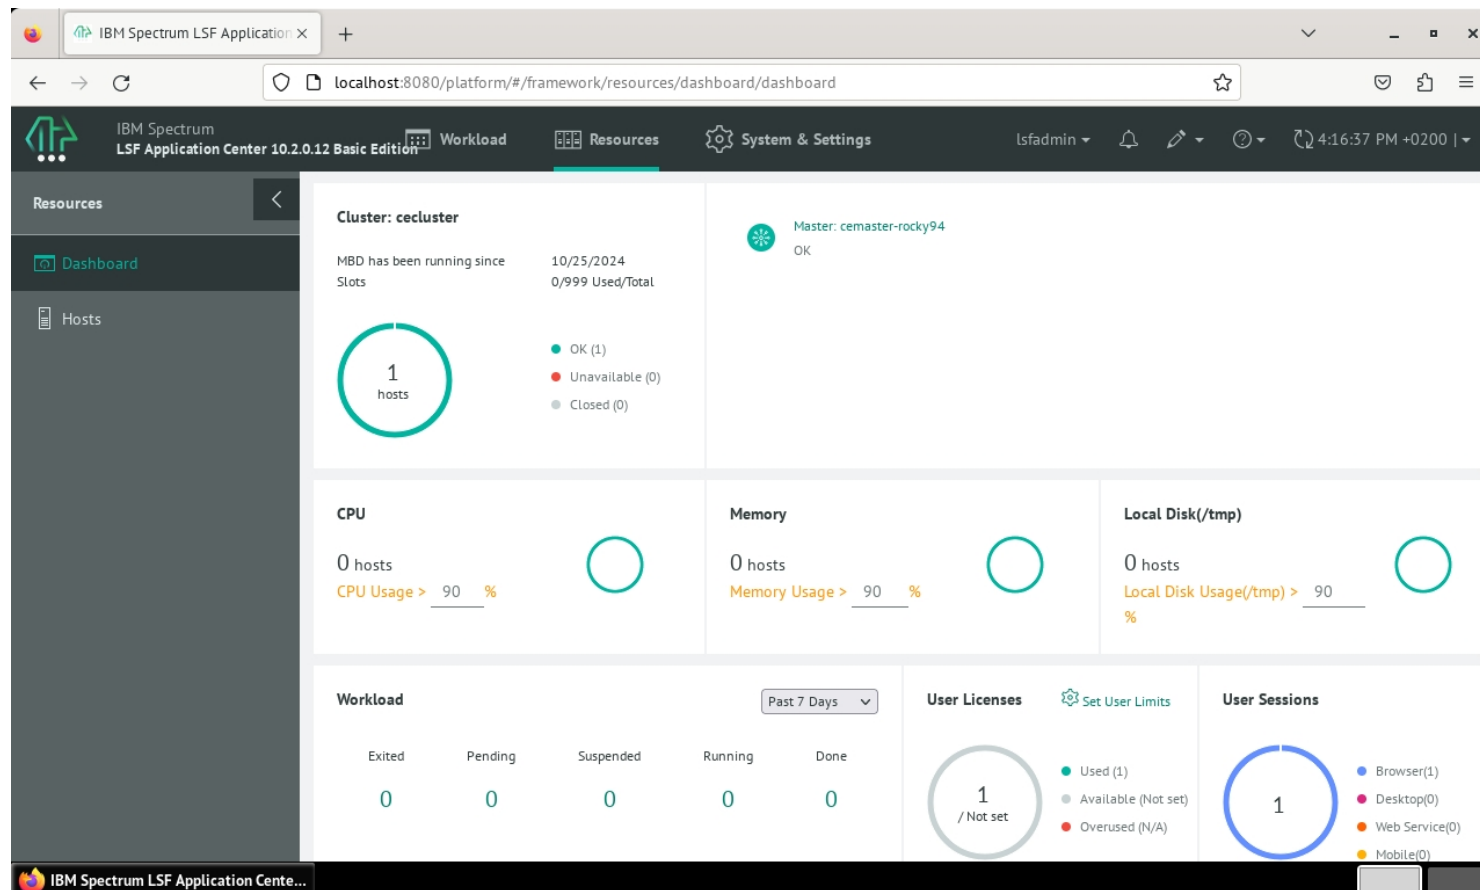
`curl -fsSL https://ibm.box.com/shared/static/h0rjy9n4yh4x6zz85sgjty01hl6f1kbl.sh | sh`



Application Center Basic Edition is part of LSF-CE



A one-node (cemaster) cluster...



Template for calculate-pi, to be executed on CodeEngine

Activities Applications Firefox Oct 25 16:18 de

IBM Spectrum LSF Application Center 10.2.0.12 — Mozilla Firefox

localhost:8080/platform/#/submissionFormDetail

Submission Form: generic

1. Job Options * 2. Notification 3. Review

Command *	calculate_pi.sh	Job Name	
CE size	0.5 vCPU 1 GB	CE image	rocky
Application	codeengine	Queue	normal
Iterations	1000	Output file	out.txt

Close Back Next Submit

Select the size of the CodeEngine instance

Activities Applications Firefox Oct 25 16:21 de

IBM Spectrum LSF Application Center 10.2.0.12 — Mozilla Firefox

localhost:8080/platform/#/submissionFormDetail

Submission Form: generic Open application menu

1. Job Options 2. Notification 3. Review

Command * calculate_pi.sh

CE size <Use Default>

Application 4 vCPU 32 GB

Iterations 6 vCPU 12 GB

6 vCPU 24 GB

6 vCPU 48 GB

8 vCPU 16 GB

8 vCPU 32 GB

10 vCPU 32 GB

Job Name

CE image rocky

Queue normal

Output file out.txt

Close Back Next Submit

LSF job 105 submitted via Application Center on CodeEngine

Jobs

Run your code to complete tasks. [Learn more](#)

Job runs

Jobs

Job name

Filter...

▼

Status

Filter...

▼

🔍

Search

⚙️

<input type="checkbox"/>	Job run name	Job name	Status	Created	↓	Duration	CPU and memory	Array size
<input type="checkbox"/>	105-jobrun-g8hug	105	<div><div></div></div> Running (1 of 1)	10/25/2024, 4:56:58 PM		00:00:19	0.5 vCPU / 1 GB	1

Result/output of LSF job 105...

The screenshot displays the IBM Spectrum LSF Application Center interface. On the left, a sidebar shows navigation options: Workload, New Workload, Workload, By Queue, By Group, and VNC Consoles. The main panel is titled 'Workload' and shows a list of jobs. Job 105 is selected, and its details are shown in the 'Data' tab. The job is named 'calculate_pi.sh' and is located at '/home/lfsadmin/gene'. The output file 'out.txt' is checked. A red circle highlights the output text in the 'out.txt' file, which reads: '37 Calculated value of Pi: 3.140592653839796'.

Queue Summary

Workload

calculate_pi.sh 105

Summary Data

Location: /home/lfsadmin/gene

Tail Open

File Name

☒ out.txt

out(1).txt

Save

Open

13

14

15 # LSBATCH: User input

16 calculate_pi.sh 1000

17

18

19 Successfully completed.

20

21 Resource usage summary:

22

23 CPU time : 2.42 sec.

24 Max Memory : 84 MB

25 Average Memory : 84.00 MB

26 Total Requested Memory : -

27 Delta Memory : -

28 Max Swap : 6626 MB

29 Max Processes : 5

30 Max Threads : 16

31 Run time : 35 sec.

32 Turnaround time : 36 sec.

33

34 The output (if any) follows:

35

36

37 Calculated value of Pi: 3.140592653839796

Plain Text Tab Width: 8 Ln 37, Col 42 INS

root@cemaster-rocky94:~

IBM Spectrum LSF Application Center

out(1).txt (~Downloads) - gedit

One step further, run LSF/SLURM master as CE instance...
(using guacamole/fvwm2/xterm), due to container restrictions...

```
root@cemaster-00001-deployment-5b6b5567b4-wq5tc: ~  
Modifying LSF configuration  
=====
```

Stopping the LSF subsystem
Starting the LSF subsystem

```
Remove any orphaned CE_jobs/jobruns  
=====
```

Creating codeengine scripts
=====

```
./install_lsf_ce_gateway.sh: 1: docker: not found  
Do you want to create new images? [y|n] <Enter> for n) n
```

```
Create /tmp/INPUT.txt  
=====
```

You may try:

```
bsub -app codeengine -I date  
bsub -app codeengine -I cat /etc/os-release  
bsub < /tmp/test.job.sh  
bsub < /tmp/testarray.sh  
bsub < /tmp/teststdout.sh  
time bsub -app codeengine -I date  
root@cemaster-00001-deployment-5b6b5567b4-wq5tc:~#
```

